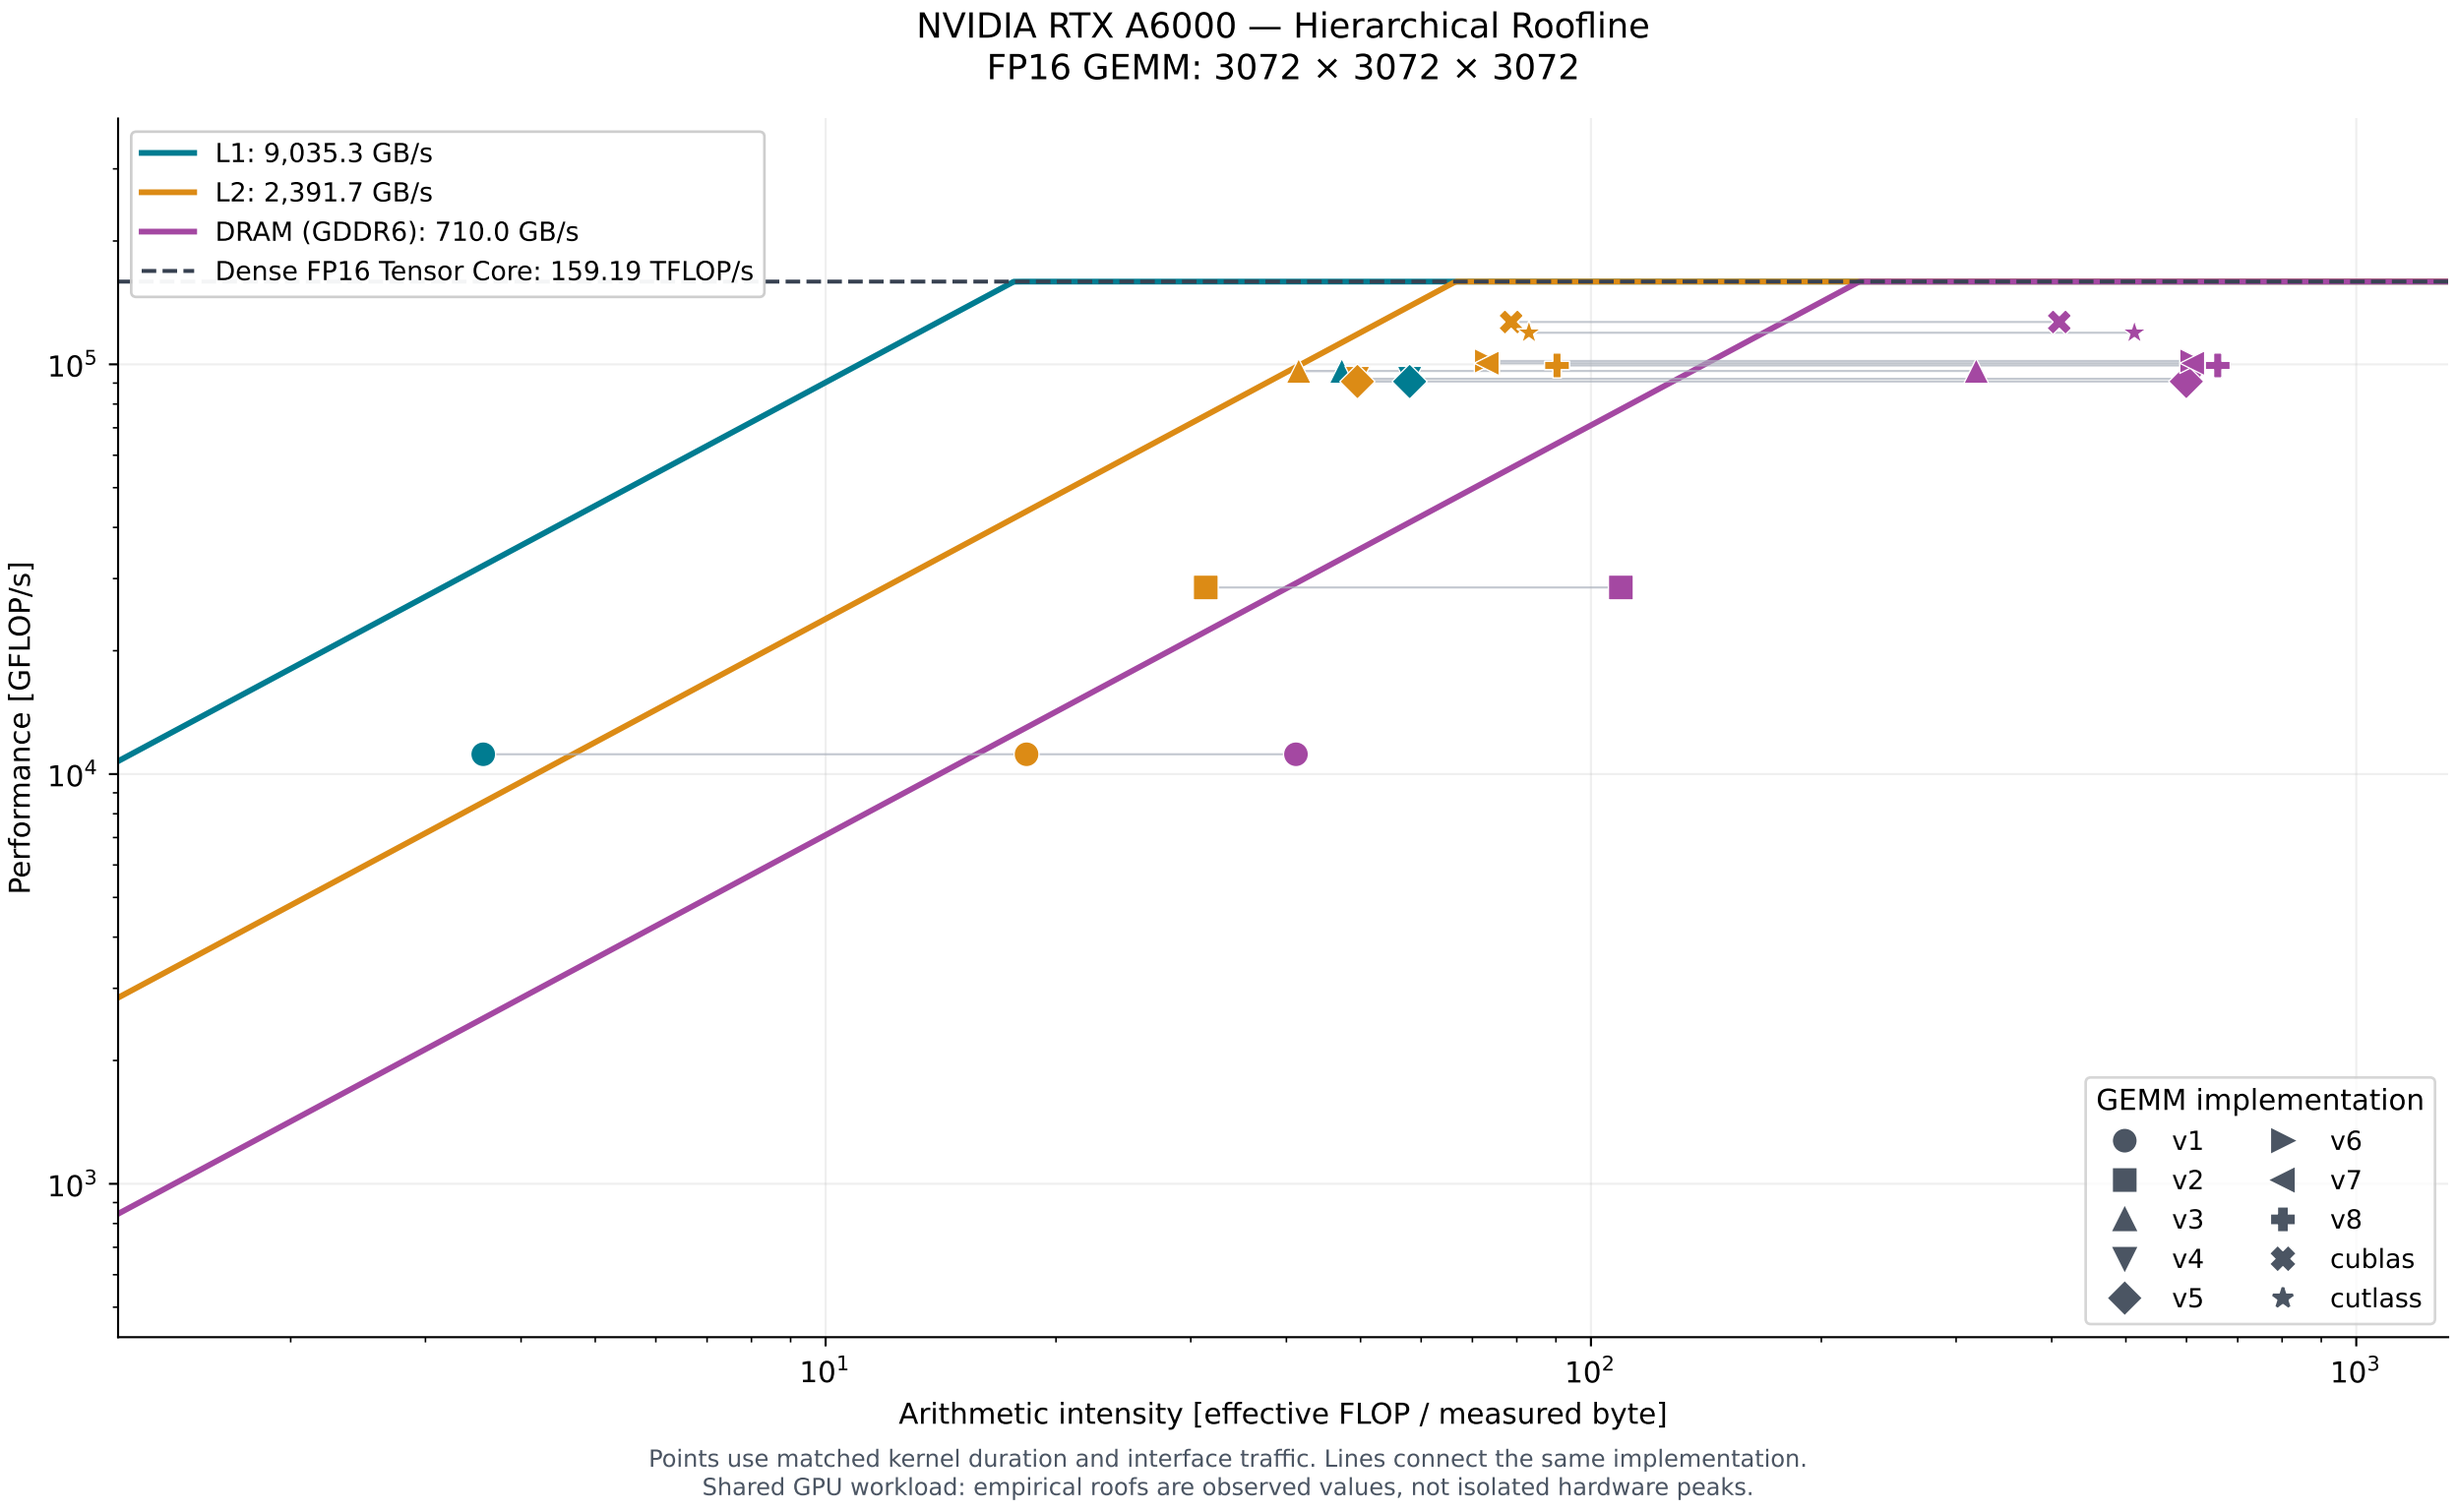
NVIDIA RTX A6000 — Hierarchical Roofline
FP16 GEMM: 3072 × 3072 × 3072
L1: 9,035.3 GB/s
L2: 2,391.7 GB/s
DRAM (GDDR6): 710.0 GB/s
Dense FP16 Tensor Core: 159.19 TFLOP/s
Performance [GFLOP/s]
Arithmetic intensity [effective FLOP / measured byte]
GEMM implementation
v1
v2
v3
v4
v5
v6
v7
v8
cublas
cutlass
Points use matched kernel duration and interface traffic. Lines connect the same implementation.
Shared GPU workload: empirical roofs are observed values, not isolated hardware peaks.